










Target: Watch
Resolution: 1920 x 1080








Reference: Adapted from an Ultralytics tutorial and documentation for the Microsoft Florence-2 model (Click to view)




visual embeddings




text + location embeddings




text + location tokens




Florence-2 Report	
Hypothesis Testing	Zero-Shot VLM Evaluation
Decision: FAIL TO REJECT H0 P-Value: 6.22e-01 T-Stat: 0.49 α : 0.05	Samples: 183 images Succ (Pt/Gr): 30.1% / 30.6% Err (Pt/Gr): 269.8 / 261.7 px Sens: 111.4px Cons: 0.92 Dual-Succ: 22.4%



Point to...
(1151, 979)



Grasp the...
(1151, 981)



Mathematical Framework

Standard Performance Metrics: Prompt sensitivity, prompt consistency, per-category error, and per-prompt error are derived from these base metrics and reported in the aggregate JSON for each model.

$$E_{pixel} = \sqrt{(x_{pred} - x_{gt})^2 + (y_{pred} - y_{gt})^2} \quad (1)$$

$$E_{norm} = \frac{E_{pixel}}{\sqrt{w^2 + h^2}} \quad (2)$$

$$S_{rate} = \frac{1}{N} \sum_{i=1}^N \mathbf{1}[E_{pixel,i} < \tau] \quad (3)$$